

**IT00CGI9-3002**

# **GPU Programming**

**Slide set #5:  
Multi-GPU Programming**

**Fall 2023, period 1**

**Jan Westerholm  
Åbo Akademi University**

# Multi-GPU Program Design and Implementation

- Many supercomputers and computer clusters consist of nodes with 2 CPU processors dies, each with tens of CPU cores, and 4 or even 8 GPUs.
- Running a GPU program on one of these nodes therefore has direct access to 4 or 8 GPUs.
  - Distributed memory programming with e.g. MPI not needed but may be helpful
- The relevant questions are now: how do we access (access = copy data to/from, run programs on) these GPUs and how do we subdivide our computational task between the GPUs?

# Multi-GPU Program Design

- Central idea: Distribute the thread blocks

```
kernel<<<threadblocks,threads_in_block>>>(...);
```

- as smaller threadblock bunches among the GPUs:

```
bunch = (threadblocks + #ofGPUs - 1) / #ofGPUs
for ( int i = 0; i < #ofGPUs; ++i )
{
    cudaSetDevice(i);
    //copy all data to device i
    cudaMemcpy(..., cudaMemcpyHostToDevice)
    int tbstart = i*bunch;
    kernel<<<bunch,threads_in_block>>>(...,tbstart);
}
```

# Multi-GPU Program Design

- Change your GPU program

```
blockIdx.x  $\Rightarrow$  blockIdx.x + tbstart
```

- Copy data back to CPU

```
for ( int i = 0; i < #ofGPUs; ++i )  
{  
    cudaSetDevice(i);  
    // copy result back to CPU and aggregate  
    cudaMemcpy(..., cudaMemcpyDeviceToHost)  
}
```

- This can also be made to work for  $\#ofGPUs = 1$ .
- Now the program can be made to run on all available GPUS, `cudaGetDeviceCount ( int* count )`

# Multi-GPU Program Design

- This design can be extended to MPI programs, where the GPU program is subdivided to work on GPUs on separate nodes (two levels: first which node, then which GPU!)
- The full data should now be subdivided among all GPUs

$(\# \text{ of nodes}) * (\# \text{ of GPUs per node})$

using the rank of the MPI program, one MPI process per GPU, as a global GPU ID.